

Applied Data Analytics II - Course Project

Diallo Youssouf

Date completed: 05/04/2019

Research questions

I will be studying customers behaviors in a Black Friday sales

1- Predict the amount purchased by a client using Regression Analysis,

2- Define which categories of Age tend to buy a lot using Classification

3- Determine which city is more productive

4- Run Cluster analysis for segmenting customers based on their spending

We argue that by answering these questions, it will help the retail store to understand their businesses and make data driven decision.

1. Importing Libraries

```
► In [1]: import pandas as pd
import numpy as np
import seaborn as sn
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from pandas.plotting import scatter_matrix
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
from matplotlib import rcParams
rcParams['figure.dpi'] = 90
sn.set_style("whitegrid")
sn.set_context("poster")
```

2. Loading dataset

The dataset that I am using is from Kaggle and it can be found at this

link <https://www.kaggle.com/mehdidag/black-friday/version/1>
(<https://www.kaggle.com/mehdidag/black-friday/version/1>)

once you click on the link, you will need to go under download button to download

```
In [2]: ## Load the csv file into a pandas Data Frame
blackFriday =pd.read_csv('Project_Diallo_Youssouf_Dataset.csv')
```

```
In [3]: ## check the number of rows and columnns of the data
blackFriday.shape
```

```
Out[3]: (537577, 12)
```

```
In [4]: ## check the data types of each features in the data set
blackFriday.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 537577 entries, 0 to 537576
Data columns (total 12 columns):
User_ID                537577 non-null int64
Product_ID             537577 non-null object
Gender                 537577 non-null object
Age                   537577 non-null object
Occupation             537577 non-null int64
City_Category          537577 non-null object
Stay_In_Current_City_Years  537577 non-null object
Marital_Status         537577 non-null int64
Product_Category_1     537577 non-null int64
Product_Category_2     370591 non-null float64
Product_Category_3     164278 non-null float64
Purchase               537577 non-null int64
dtypes: float64(2), int64(5), object(5)
memory usage: 49.2+ MB
```

3. Data set Description

The dataset contains 537,577 observations about black Friday sales a for a retail company.

It contains 12 columns with numerical and categorical data types

4. Columns or Attributes

USER_ID : Define the User id, which will help identify a user who made a purchase from the store

Product_ID : Help identify a prodcut in the store

Gender: Define the user gender which will be either Male (M) or Female(F)

Age : Define the age of user, which is given in a range of values.

Occupation : Define the ID occupation of each customer

City_Category : Define the category of each city (A, B, and C)

Stay_In_Current_City_Years : Define how many years, a customer resides in that city

Marital_Status : Define the customer's status, which is either married coded by 1 or not married coded by 0

The retail company is selling three categories of products

Product_Category_1 : Define the number of items purchased within the Product_Category_1

Product_Category_2 : Define the number of items purchased within the Product_Category_2

Product_Category_3 : Define the number of items purchased within the Product_Category_3

Purchase : Define the amount spent by each customer to purchase the three products. it is in US dollars

5. Data Exploration

► In [6]:

```
## This print the first 5 elements
blackFriday.head()
```

Out[6]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	1000001	P00069042	F	0-17	10	A	2	0
1	1000001	P00248942	F	0-17	10	A	2	0
2	1000001	P00087842	F	0-17	10	A	2	0
3	1000001	P00085442	F	0-17	10	A	2	0
4	1000002	P00285442	M	55+	16	C	4+	0

► In [7]: `## This print the last 5 elements`
`blackFriday.tail()`

Out[7]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years
537572	1004737	P00193542	M	36-45	16	C	1
537573	1004737	P00111142	M	36-45	16	C	1
537574	1004737	P00345942	M	36-45	16	C	1
537575	1004737	P00285842	M	36-45	16	C	1
537576	1004737	P00118242	M	36-45	16	C	1

► In [8]: `## Let check if we have null`
`blackFriday.isnull().values.any()`

Out[8]: True

yes we have null values in the dataset, but I will handle it later

To better understand our data let's find some statistics

what was the most profitable product on Black Friday ?

► In [8]: `## Let's check products that are sold the most and just pick the top 5`
`blackFriday.groupby(["Product_ID"]).sum().sort_values("Purchase", ascending=False)`

Out[8]:

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Proc
Product_ID						
P00025442	1590774903		13112	635	1586	3172.0
P00110742	1595784075		12933	637	1591	3182.0
P00255842	1358323490		11196	526	21664	0.0
P00184942	1428158481		11680	589	1424	11392.0
P00059442	1388200080		11456	567	8304	11072.0

We can see from the result above the product with ID (P00025442) was the most sold with a total purchase of \$ 27,532426

what was the most profitable City on Black Friday ?

```

In [9]: group = blackFriday.groupby(["City_Category"]).sum()
total_price = group["Purchase"].groupby(level=0, group_keys=False)
total_price.nlargest(5)

```

```

Out[9]: City_Category
A      1295668797
B      2083431612
C      1638567969
Name: Purchase, dtype: int64

```

from the result above the city category B was the most profitable on Black Friday

the purchase amount was \$ 2,083 431 612 something above 2 Billion dollars

Determine which group of people buy a lot

```

In [10]: ## this count the number in each category
blackFriday.Gender.value_counts()

```

```

Out[10]: M      405380
F       132197
Name: Gender, dtype: int64

```

The result above shows that the retail company recorded 405380 transactions from Males

which represent 75.40 % and 132197 transactions which also represent 24.6 % from Females

```

In [11]: ## this line of code group by Marital status
blackFriday.groupby(["Marital_Status"]).sum().sort_values("Purchase", ascending=False)

```

```

Out[11]:

```

	User_ID	Occupation	Product_Category_1	Product_Category_2	Product_Category_3
Marital_Status					
0	318759372904	2526251	1662649	2154478.0	1237
1	220425975246	1818828	1184115	1492932.0	843

```

In [12]: ## we check the number of people who are married or unmarried
blackFriday.Marital_Status.value_counts()

```

```

Out[12]: 0    317817
         1    219760
         Name: Marital_Status, dtype: int64

```

This result shows that 59 % of the customers are married and 41 % are unmarried

```

In [80]: ## Let group the data set by Purchase amount
#blackFriday.groupby('Purchase').sum()

```

This result above shows us that the minimum amount purchase was 185 dollars and the maximum was 23,961 dollars

Let's run some descriptive statistics

```

In [9]: ## this line of code compute the descriptive statistics
blackFriday.describe()

```

```

Out[9]:

```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Produ
count	5.375770e+05	537577.00000	537577.000000	537577.000000	370591.000000	
mean	1.002992e+06	8.08271	0.408797	5.295546	9.842144	
std	1.714393e+03	6.52412	0.491612	3.750701	5.087259	
min	1.000001e+06	0.00000	0.000000	1.000000	2.000000	
25%	1.001495e+06	2.00000	0.000000	1.000000	5.000000	
50%	1.003031e+06	7.00000	0.000000	5.000000	9.000000	
75%	1.004417e+06	14.00000	1.000000	8.000000	15.000000	
max	1.006040e+06	20.00000	1.000000	18.000000	18.000000	

This result gives us a brief description of mean, standard deviation, max, min and etc of the data set

```
► In [10]: ## Let's calculate the total purchase and group by Gender and Marital_Status  
blackFriday.groupby(["Gender", "Marital_Status"]).sum().sort_values("Purchase", asc
```

```
Out[10]:
```

			User_ID	Occupation	Product_Category_1	Product_Category_2	Produ
	Gender	Marital_Status					
	M	0	241555332289	2042528	1235727	1630874.0	
		1	165024843194	1411190	871336	1117320.0	
	F	0	77204040615	483723	426922	523604.0	
		1	55401132052	407638	312779	375612.0	

```
► In [11]: ## compute the number of product category 1 that was sold  
blackFriday.Product_Category_1.sum()
```

```
Out[11]: 2846764
```

```
► In [12]: ## compute the number of product category 2 that was sold  
blackFriday.Product_Category_2.sum()
```

```
Out[12]: 3647410.0
```

```
► In [13]: ## compute the number of product category 3 that was sold  
blackFriday.Product_Category_3.sum()
```

```
Out[13]: 2081376.0
```

The conclusion that we can draw from this is that product category 3 was the most sold

3647410 items sold

6. Preprocessing

```
► In [14]: ## Let's do some data cleaning  
blackFriday.shape
```

```
Out[14]: (537577, 12)
```

```
► In [4]: ## check data types  
blackFriday.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 537577 entries, 0 to 537576  
Data columns (total 12 columns):  
User_ID                537577 non-null int64  
Product_ID            537577 non-null object  
Gender                537577 non-null object  
Age                  537577 non-null object  
Occupation            537577 non-null int64  
City_Category         537577 non-null object  
Stay_In_Current_City_Years  537577 non-null object  
Marital_Status        537577 non-null int64  
Product_Category_1    537577 non-null int64  
Product_Category_2    370591 non-null float64  
Product_Category_3    164278 non-null float64  
Purchase              537577 non-null int64  
dtypes: float64(2), int64(5), object(5)  
memory usage: 49.2+ MB
```

From the result above, we can see that Product_Category_2 and Product_Category_3 have NaN values

Instead of deleting the attributes Product_Category_2 and Product_Category_3, we will replace the NaN values by zero

```
► In [3]: ## this line of code replaces the NaN values with zero  
blackFriday=blackFriday.fillna(0)
```

```
► In [4]: ## Let's check again if we still have NaN  
blackFriday.isnull().values.any()
```

Out[4]: False

we can see that we do not have any more missing

Let's do some data types transformations on Product_Category_2 and Product_Category_3

```
► In [5]: ## we replace the data types by int, we want to keep everything in int instead of f  
blackFriday['Product_Category_2'] = blackFriday.Product_Category_2.astype(np.int64)  
blackFriday['Product_Category_3'] = blackFriday.Product_Category_3.astype(np.int64)
```


► In [6]: *## we check again to see if there is change on the data types*
 blackFriday.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 537577 entries, 0 to 537576
Data columns (total 12 columns):
User_ID                537577 non-null int64
Product_ID            537577 non-null object
Gender                537577 non-null object
Age                   537577 non-null object
Occupation            537577 non-null int64
City_Category         537577 non-null object
Stay_In_Current_City_Years  537577 non-null object
Marital_Status        537577 non-null int64
Product_Category_1    537577 non-null int64
Product_Category_2    537577 non-null int64
Product_Category_3    537577 non-null int64
Purchase              537577 non-null int64
dtypes: int64(7), object(5)
memory usage: 49.2+ MB
```

we need also to transform the type of Stay_In_Current_City_Years from object to int

► In [7]: *## we need to replace the + sign with empty string and change the type to int64*
 blackFriday['Stay_In_Current_City_Years'] = blackFriday.Stay_In_Current_City_Years

Let's group the amount purchase by category Age

► In [8]: blackFriday.groupby(["Age"]).sum().sort_values("Purchase", ascending=False)

Out[8]:

	User_ID	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	P
Age						
26-35	215350137743	1696554	406963	84166	1120056	
36-45	107824726995	951060	203189	42507	579163	
18-25	97904093780	657774	177953	20641	488498	
46-50	44666364150	379645	78462	32194	250663	
51-55	37728982395	331396	66964	26979	212529	
55+	20964683314	199372	39865	13273	123256	
0-17	14746359773	129278	26206	0	72599	

This result shows that people with Age 26-35 made a lot purchase on Black Friday followed by 36-45

But, people with age range between 0-17 made the least purchase

```
► In [9]: ## Let's see the result of the datatype  
blackFriday.info()
```


```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 537577 entries, 0 to 537576  
Data columns (total 12 columns):  
User_ID                537577 non-null int64  
Product_ID            537577 non-null object  
Gender                 537577 non-null object  
Age                   537577 non-null object  
Occupation             537577 non-null int64  
City_Category          537577 non-null object  
Stay_In_Current_City_Years  537577 non-null int64  
Marital_Status         537577 non-null int64  
Product_Category_1     537577 non-null int64  
Product_Category_2     537577 non-null int64  
Product_Category_3     537577 non-null int64  
Purchase              537577 non-null int64  
dtypes: int64(8), object(4)  
memory usage: 49.2+ MB
```

Now our data is clean and ready for building model

7. Visualization

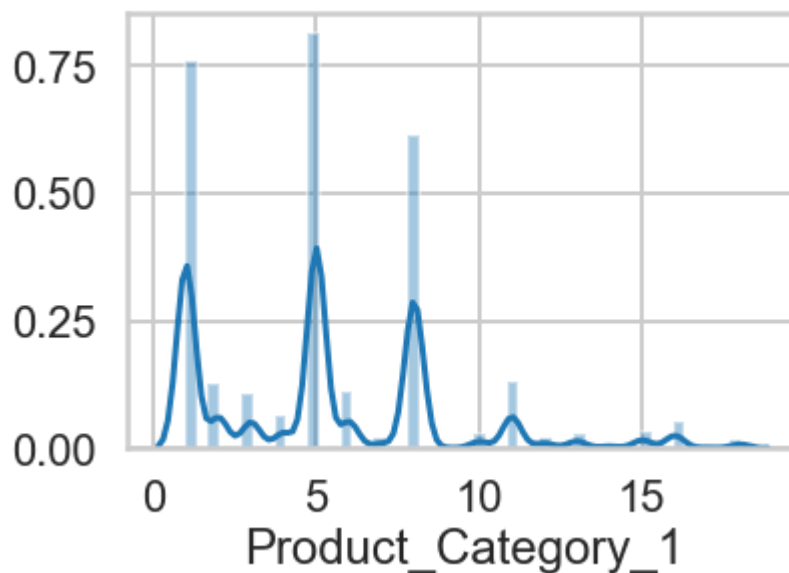
```
In [28]: ## check the distribution of purchase  
sn.distplot(blackFriday.Purchase)
```

Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x2a9042443c8>



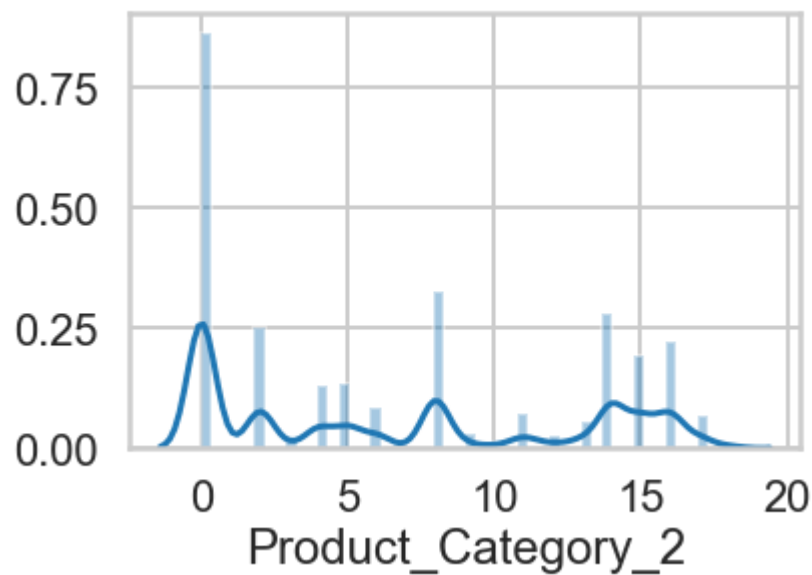
```
In [29]: ## check the distribution of Product Category 1  
sn.distplot(blackFriday.Product_Category_1)
```

Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x2a904232048>




```
► In [30]: ## check the distribution of Product category 2  
sn.distplot(blackFriday.Product_Category_2)
```

Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x2a903b96c18>



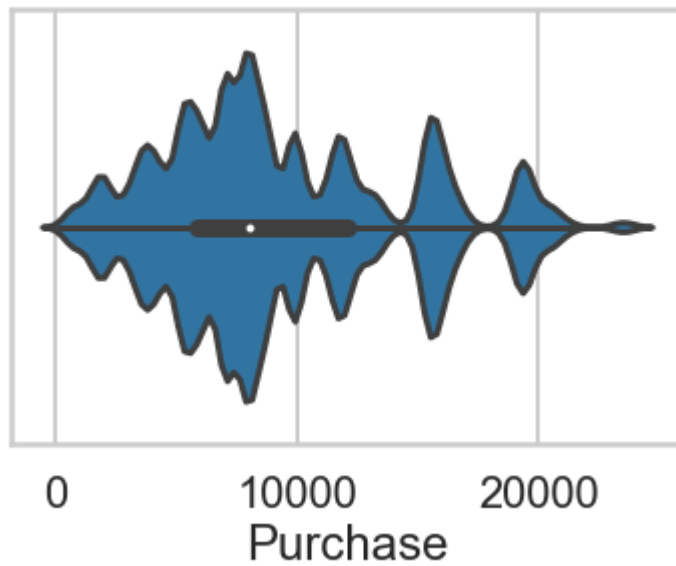
```
► In [31]: ## check the distribution of Product category 3  
sn.distplot(blackFriday.Product_Category_3)
```

Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x1f41ebc5668>



```
► In [32]: ## Let find the distribution and the density  
sn.violinplot(x = "Purchase", data=blackFriday)
```

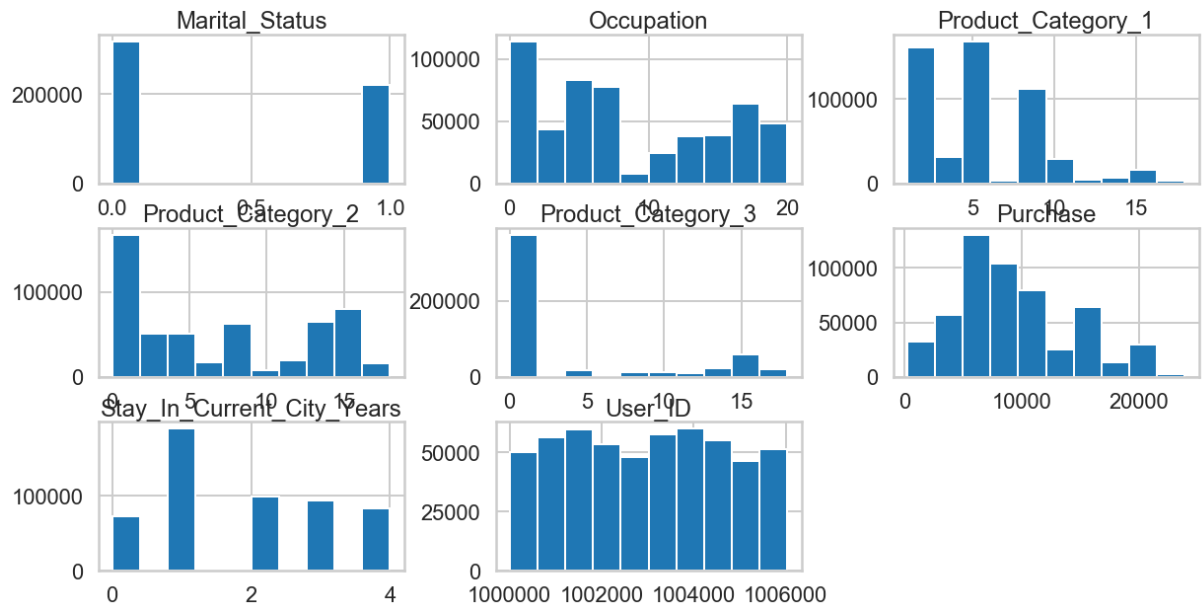
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x1f41ec65e80>



The violinplot is used to visualize the distribution of the data and its probability density

► In [33]: *## check the histogram of the data set for the numerical values*
 blackFriday.hist(figsize=(20,10))

Out[33]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000001F41F0DFCC0>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F41ED025C0>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F41ED28E80>],
 [<matplotlib.axes._subplots.AxesSubplot object at 0x000001F41ED58828>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F41ED881D0>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F41EDABB38>],
 [<matplotlib.axes._subplots.AxesSubplot object at 0x000001F41EDDB4E0>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F41EE01E80>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F41EE01EB8>]],
 dtype=object)




the result above gives us more visibility of distributions

```
In [35]: plt.scatter(blackFriday.Product_Category_1, blackFriday.Purchase)
plt.xlabel("Product Category 1")
plt.ylabel("Amount Purchase")
plt.title("Relationship between Product Category 1 and Amount Purchase")
```

Out[35]: Text(0.5, 1.0, 'Relationship between Product Category 1 and Amount Purchase')

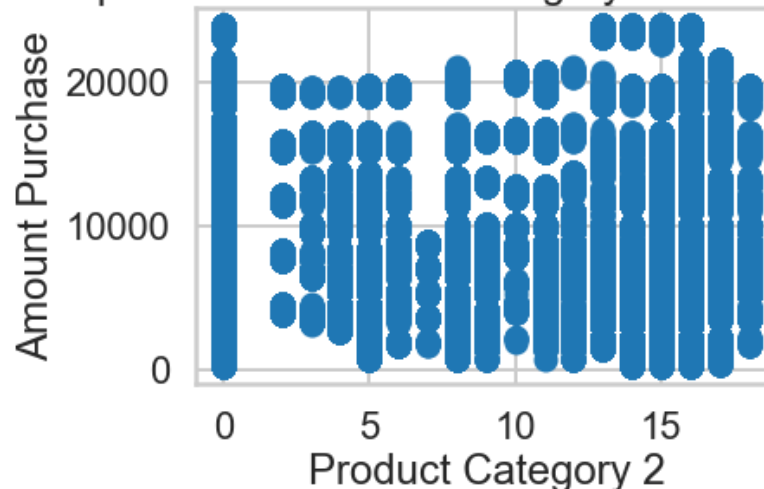
Relationship between Product Category 1 and Amount Purchase



```
In [36]: plt.scatter(blackFriday.Product_Category_2, blackFriday.Purchase)
plt.xlabel("Product Category 2")
plt.ylabel("Amount Purchase")
plt.title("Relationship between Product Category 2 and Amount Purchase")
```

Out[36]: Text(0.5, 1.0, 'Relationship between Product Category 2 and Amount Purchase')

Relationship between Product Category 2 and Amount Purchase




```

In [37]: plt.scatter(blackFriday.Product_Category_3, blackFriday.Purchase)
plt.xlabel("Product Category 3")
plt.ylabel("Amount Purchase")
plt.title("Relationship between Product Category 3 and Amount Purchase")

```

Out[37]: Text(0.5, 1.0, 'Relationship between Product Category 3 and Amount Purchase')

Relationship between Product Category 3 and Amount Purchase

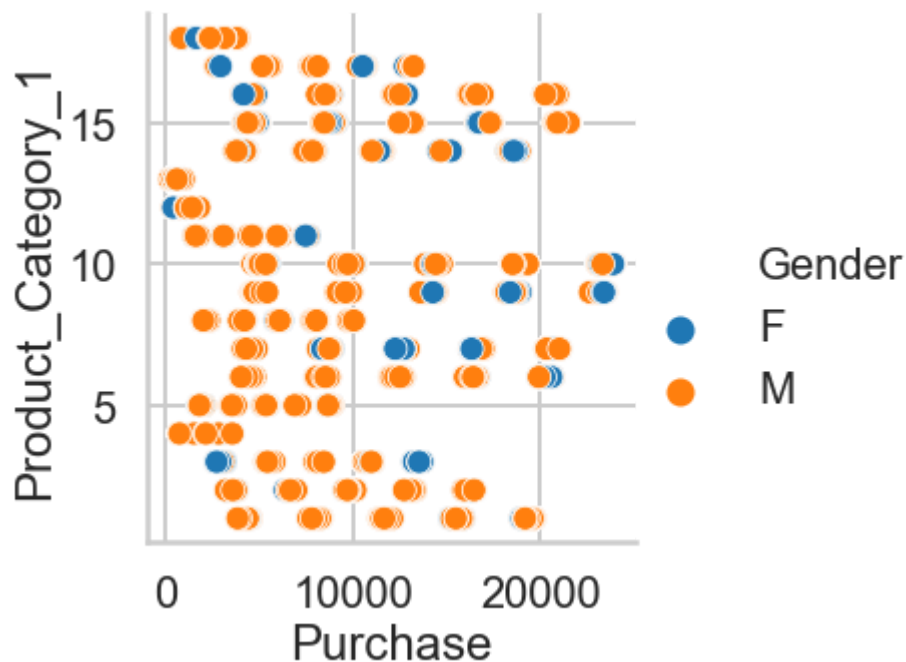


```

In [34]: ## check the relationship between Purchase and Product category based on gender
sn.relplot(x="Purchase", y="Product_Category_1", hue="Gender", data=blackFriday)

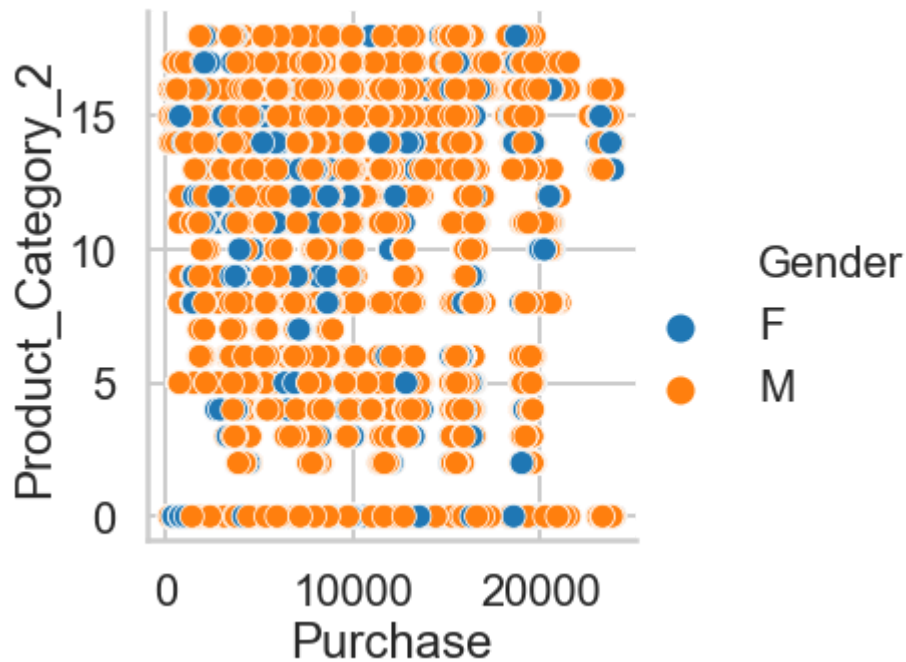
```

Out[34]: <seaborn.axisgrid.FacetGrid at 0x1f41efa09b0>



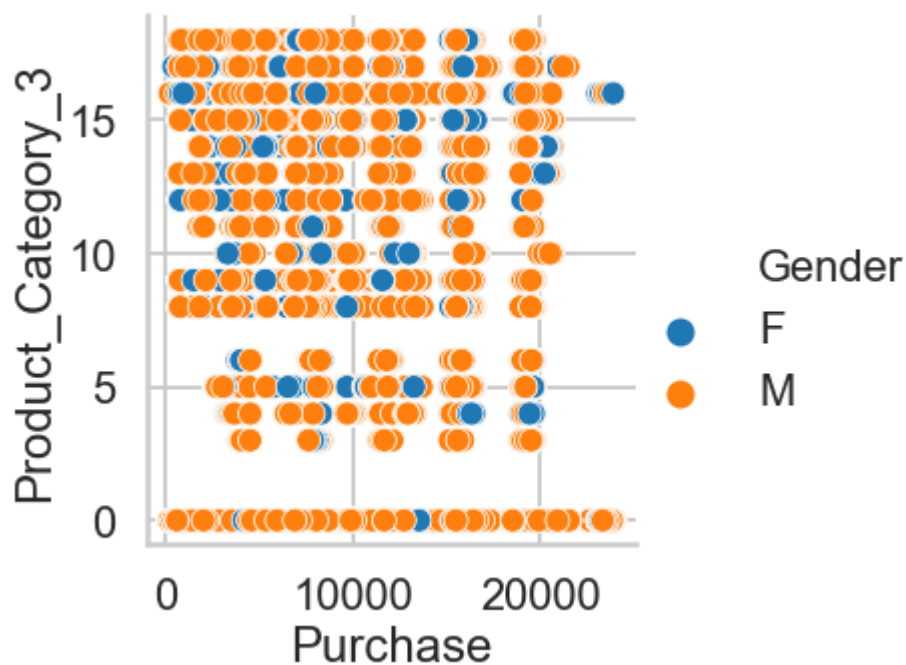

```
In [38]: ## check the relationship between Purchase and Product category 2 based on gender  
sn.relplot(x="Purchase", y="Product_Category_2", hue="Gender", data=blackFriday)
```

Out[38]: <seaborn.axisgrid.FacetGrid at 0x1f41f174d68>




```
In [39]: ## check the relationship between Purchase and Product category 3 based on gender  
sn.relplot(x="Purchase", y="Product_Category_3", hue="Gender", data=blackFriday)
```

Out[39]: <seaborn.axisgrid.FacetGrid at 0x1f41f1f5a90>



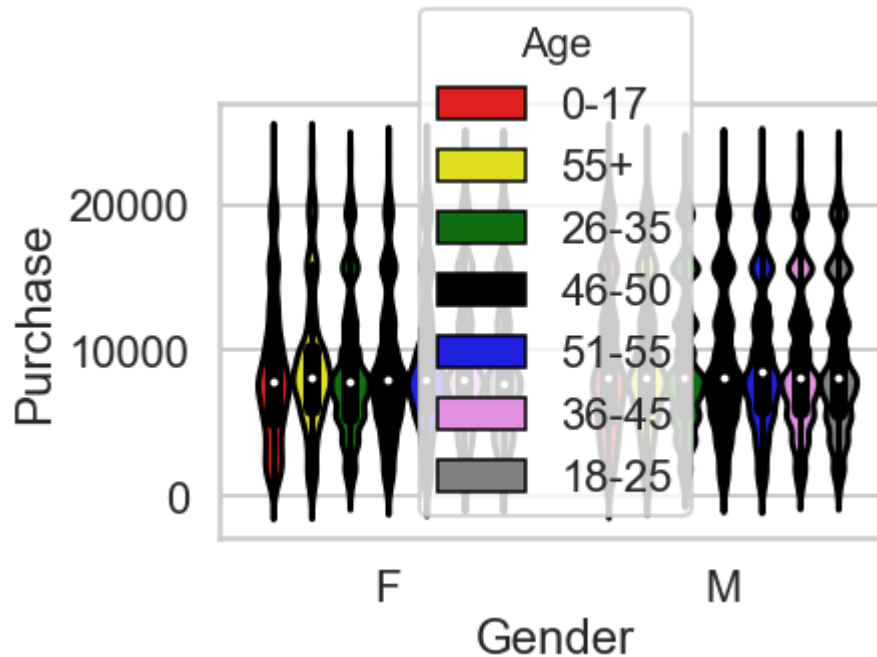
```
In [40]: sn.violinplot(x="Gender", y="Product_Category_1",hue='City_Category', data=blackFr
```

```
Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x1f41f2586d8>
```



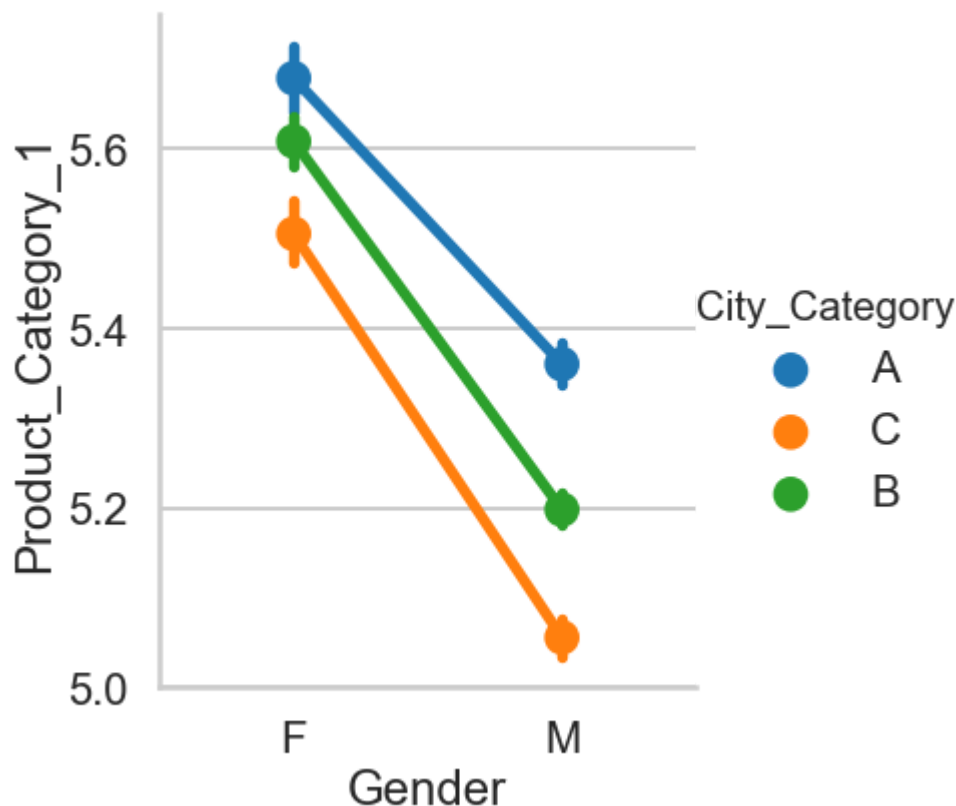
```
In [41]: sn.violinplot(x="Gender", y="Purchase",hue='Age', data=blackFriday,
                palette=["red", "yellow", "green", "black", "blue", "violet", "gray"],fiz
```

```
Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x1f41f2e1dd8>
```



► In [42]: `## check the relationship between numerical and categorical variables`

```
sn.catplot(x="Gender", y="Product_Category_1", hue="City_Category", data=blackFriday
```



► In [12]: `#Let's verify the data types`

```
blackFriday.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 537577 entries, 0 to 537576
Data columns (total 12 columns):
User_ID                537577 non-null int64
Product_ID             537577 non-null object
Gender                 537577 non-null object
Age                   537577 non-null object
Occupation             537577 non-null int64
City_Category          537577 non-null object
Stay_In_Current_City_Years  537577 non-null int64
Marital_Status         537577 non-null int64
Product_Category_1     537577 non-null int64
Product_Category_2     537577 non-null int64
Product_Category_3     537577 non-null int64
Purchase               537577 non-null int64
dtypes: int64(8), object(4)
memory usage: 49.2+ MB
```

I will be working with the entire dataset 537577 records

I will reserve 10 % for validation

```
► In [11]: ## This line of code takes 10 % for validation  
blackFridayValidation=blackFriday.sample(frac=0.1, random_state=1)
```

```
► In [12]: ## this is the validation data set  
blackFridayValidation.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 53758 entries, 94689 to 449041  
Data columns (total 12 columns):  
User_ID                53758 non-null int64  
Product_ID            53758 non-null object  
Gender                 53758 non-null object  
Age                   53758 non-null object  
Occupation             53758 non-null int64  
City_Category          53758 non-null object  
Stay_In_Current_City_Years  53758 non-null int64  
Marital_Status         53758 non-null int64  
Product_Category_1     53758 non-null int64  
Product_Category_2     53758 non-null int64  
Product_Category_3     53758 non-null int64  
Purchase               53758 non-null int64  
dtypes: int64(8), object(4)  
memory usage: 5.3+ MB
```

Our validation data has 53,758 records

Let's copy the data set for Regression, Classification and Clustering

```
► In [14]: blackFridayRegression=blackFriday      ### use for the regression  
blackFridayClassification=blackFriday          ### use for classification  
blackFridayClustering=blackFriday              ### use for clustering
```

```
► In [14]: # this check for duplicated  
blackFriday.duplicated(keep=False).value_counts()
```

```
Out[14]: False    537577  
dtype: int64
```

```
► In [ ]: sn.pairplot(blackFriday, kind="reg");
```

```
► In [ ]: corr = blackFriday.corr()  
sn.heatmap(corr)
```

this result shows us that we do not have duplicate values in our data set

Algorithm #1 Linear Regression

Predict the purchase amount

Y = Purchase (called target data in python, and referred to as the dependent variable or response variable)

X = independent variables, or explanatory variables (Occupation, Marital_Status, Stay_In_Current_City_Years, Product_category1, 2 and 3)

we will use to fit a linear regression model and predict amount Purchase. We will use the r squared method to estimate the accuracy.

► In [31]: `## we check again to see the data
blackFridayRegression.head()`

Out[31]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Mari
0	1000001	P00069042	F	0-17	10	A		2
1	1000001	P00248942	F	0-17	10	A		2
2	1000001	P00087842	F	0-17	10	A		2
3	1000001	P00085442	F	0-17	10	A		2
4	1000002	P00285442	M	55+	16	C		4

Fitting Linear Regression using sklearn

► In [29]: `## we import the libraries for building the linear model
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score`

► In [30]: `## creates a list of independent variables
predictors = ['Occupation', 'Stay_In_Current_City_Years', 'Marital_Status',
 'Product_Category_1', 'Product_Category_2', 'Product_Category_3']`

we create our target and independent variables

► In [32]: `X=blackFridayRegression[predictors] ## this hold the independent variable
Y=blackFridayRegression.Purchase ## hold the target variable`

Splitting the Data

```
► In [33]: # we need to split the data into training set and test set
# we do 70 % training and 30 % test
X_train, X_test, y_train, y_test=train_test_split(X,Y,test_size=0.3,random_state=1
```

Building Regression model

```
► In [34]: lm = LinearRegression()
lm
```

Out[34]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

Fit the Regression model

```
► In [35]: lm.fit(X_train,y_train)
```

Out[35]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

Make prediction using testing set

```
► In [36]: model_yPred=lm.predict(X_test)
```

```
► In [37]: # The coefficients
print('Coefficients: \n', lm.coef_)
# The mean squared error
print("Mean squared error: %.2f"
      % mean_squared_error(y_test, model_yPred))
# Explained variance score: 1 is perfect prediction
print('R squared : %.2f' % r2_score(y_test, model_yPred))
```

```
Coefficients:
[ 13.46204589   6.46363812  36.6513731 -319.25980883   9.23250636
 151.0691684 ]
Mean squared error: 21561218.90
R squared : 0.13
```

The result above shows the mean squared error and the coefficient

the mean squared error is very high which means the model did not perform well

```
► In [38]: lm.score(X_test, y_test)
```

```
Out[38]: 0.1291209162422322
```

the model is very terrible only 13 % of purchase amount can be explained by the chosen independent variables

Make validation using validation set

```
► In [57]: valid=blackFridayValidation.Purchase  
Xvalidation=blackFridayValidation[predictors]
```

```
► In [61]: X_train, X_valid, y_train, y_valid=train_test_split(Xvalidation,valid,test_size=0.
```

```
► In [62]: ## we predict with the validation set  
model_yPred=lm.predict(X_valid)
```

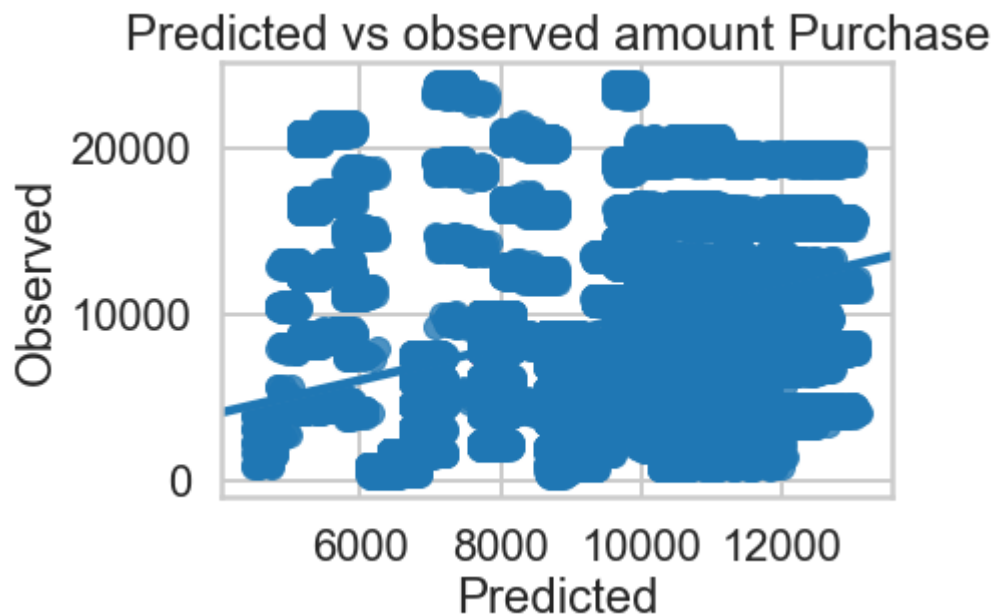
```
► In [65]: ## check the score based on the validation set  
lm.score(X_valid, y_valid)
```

```
Out[65]: 0.1333709930717688
```


Using the validation set, we find out that the accuracy score is slightly high

Regression plot

```
► In [57]: # plot relationship between observed and predicted amount purchase
sn.regplot(x=model_yPred, y=y_test)
plt.xlabel('Predicted')
plt.ylabel('Observed')
plt.title(('Predicted vs observed amount Purchase'));
```



```
► In [173]: # plot histogram of predicted amount purchase
plt.hist(lm.predict(X_test), bins=20)
plt.xlabel("Predicted amount Purchase")
plt.ylabel("Frequency")
plt.title("Distribution of predicted amount Purchase");
```





```

In [172]: # plot histogram of observed amount purchase
plt.hist(blackFridayRegression.Purchase, bins=20)
plt.xlabel("Observed amount Purchase")
plt.ylabel("Frequency")
plt.title("Distribution of observed amount Purchase")

```

Out[172]: Text(0.5, 1.0, 'Distribution of observed amount Purchase')



Let's use another technique to see if we can improve the model

We will use XGboost for Regression

```

In [32]: ## import the libraries for XGBoost
import xgboost
from sklearn.metrics import explained_variance_score

```

Splitting Data into training and testing 80 % and 20 %

```

In [190]: X_train, X_test, y_train, y_test = train_test_split(X, Y ,test_size=0.2)

```

Building XGBoost Model

```

In [191]: # Let's try XGboost algorithm to see if we can get better results
xgb = xgboost.XGBRegressor(n_estimators=100, learning_rate=0.08, gamma=0, subsamp
                           colsample_bytree=1, max_depth=7)

```

```
► In [197]: traindf, testdf = train_test_split(X_train, test_size = 0.2)
           xgb.fit(X_train,y_train)
```

```
Out[197]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bytree=1, gamma=0, importance_type='gain',
                        learning_rate=0.08, max_delta_step=0, max_depth=7,
                        min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
                        nthread=None, objective='reg:linear', random_state=0, reg_alpha=0,
                        reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
                        subsample=0.75)
```


Predict XGBoost Model

```
► In [198]: predictions = xgb.predict(X_test)
           print(explained_variance_score(predictions,y_test))
```

```
0.45930676102949164
```

We clearly see that XGboost has improved our model and the accuracy moves from 0.13 to 0.459

```
► In [199]: # plot relationship between observed and predicted amount purchase
           sn.regplot(x=predictions, y=y_test)
           plt.xlabel('Predicted')
           plt.ylabel('Observed')
           plt.title(('Predicted vs observed amount Purchase'));
```



we see that the data points are clustered around the regression line

this is the best we can get

Algorithm #2 Decision Tree Classifier

We are using classification to classify the age and the city

```
▶ In [67]: blackFridayClassification.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 537577 entries, 0 to 537576
Data columns (total 12 columns):
User_ID                537577 non-null int64
Product_ID             537577 non-null object
Gender                 537577 non-null object
Age                   537577 non-null object
Occupation             537577 non-null int64
City_Category          537577 non-null object
Stay_In_Current_City_Years  537577 non-null int64
Marital_Status         537577 non-null int64
Product_Category_1     537577 non-null int64
Product_Category_2     537577 non-null int64
Product_Category_3     537577 non-null int64
Purchase              537577 non-null int64
dtypes: int64(8), object(4)
memory usage: 49.2+ MB
```

We are going to classify the age group

```
▶ In [68]: target=blackFridayClassification.Age ## this the target
features = ['Occupation', 'Stay_In_Current_City_Years', 'Marital_Status', 'Purchase',
            'Product_Category_1', 'Product_Category_2', 'Product_Category_3']
```

```
▶ In [69]: X=blackFridayClassification[features]
```

Splitting the Data

```
▶ In [70]: # we need to split the data into training set and test set
# we do 70 % training and 30 % test
X_train, X_test, y_train, y_test=train_test_split(X,target,test_size=0.3,random_st
```

Building Decision Tree Model

```
▶ In [74]: ## import libraries for building a decision tree classifier
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics
```

```
► In [75]: # we create an object decision tree
decisionTree =DecisionTreeClassifier()
```

```
► In [76]: # We train the decision tree classifier
decisionTree.fit(X_train,y_train)
```

```
Out[76]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

Make prediction using testing set

```
► In [77]: # we predict the response for the test
y_predict=decisionTree.predict(X_test)
```

Evaluation the decision tree classifier

```
► In [78]: # Accuracy is calculted by using the test score by comparing the test values and p
print("Accuracy",metrics.accuracy_score(y_test,y_predict))
```

Accuracy 0.416837183923013

we find that the accuracy of the decision tree classifier is 41.68.2 %

vusialization decision trees libraries

```
► In [45]: ## this is the libraries for plotting
from sklearn.tree import export_graphviz
from sklearn.externals.six import StringIO
from IPython.display import Image
import pydotplus
```

Optimizing Decision tree classifier

```
► In [87]: # Let see if we can optimize the decison tree by specifying the criterion and max
decisionOptimize =DecisionTreeClassifier(criterion="entropy",max_depth=8)
```

```
► In [88]: ## we train the model
decisionOptimize.fit(X_train,y_train)
```

```
Out[88]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=8,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

```
► In [89]: # predict the response for the dataset
y_preditOpt=decisionOptimize.predict(X_test)
```

```
► In [90]: # Measure de accuracy
print("Accuracy :",metrics.accuracy_score(y_test,y_preditOpt))
```

Accuracy : 0.5068392921363643

the accuracy of the model is now 50.6 % , a bit of improvement

Visualizing the tree

```
► In [ ]: dot_data = StringIO()
export_graphviz(decisionTree, out_file=dot_data,
                filled=True, rounded=True,
                special_characters=True,feature_names = features,class_names=['0-17','18-2
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png('age.png')
Image(graph.create_png())
```

Now, we are going to classify the city

```
► In [91]: mytarget=blackFridayClassification.City_Category ## this the target
features =['Occupation','Stay_In_Current_City_Years','Marital_Status','Purchase',
           'Product_Category_1','Product_Category_2','Product_Category_3']
```

```
► In [92]: X=blackFridayClassification[features]
```

► In [93]: `blackFridayClassification.head()`

Out[93]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Mari
0	1000001	P00069042	F	0-17	10	A		2
1	1000001	P00248942	F	0-17	10	A		2
2	1000001	P00087842	F	0-17	10	A		2
3	1000001	P00085442	F	0-17	10	A		2
4	1000002	P00285442	M	55+	16	C		4

Splitting Data

► In [94]: `# we need to split the data into training set and test set`
`# we do 70 % training and 30 % test`
`X_train, X_test, y_train, y_test=train_test_split(X,mytarget,test_size=0.3,random_`

► In [95]: `# we create an object decision tree`
`decisionTree =DecisionTreeClassifier()`

► In [96]: `# We train the decision tree classifier`
`decisionTree.fit(X_train,y_train)`

Out[96]: `DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None, max_features=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, presort=False, random_state=None, splitter='best')`

Make prediction using testing set

► In [97]: `# we predict the response for the test`
`y_predict=decisionTree.predict(X_test)`

Evaluation of model

► In [98]: `# Accuracy is calculted by using the test score by comparing the test values and p`
`print("Accuracy",metrics.accuracy_score(y_test,y_predict))`

Accuracy 0.4234656547242581

the accuracy of the model is 42.23 %

We are going to use XGboost Classifier to see if we can improve the model

```

In [99]: # this import the library for encoding categorical variables
from sklearn.preprocessing import LabelEncoder

```

```

In [100]: ## Because XGBoost takes only numerical values, we are going to transform city from
number = LabelEncoder()
blackFridayClassification2 = blackFridayClassification
blackFridayClassification2['City_Category'] = number.fit_transform(blackFridayClassification2['City_Category'])

```

```

In [101]: ## import libraries for XGBoost classifier
import xgboost as xgb
from xgboost import XGBClassifier

```

```

In [102]: mytarget = blackFridayClassification2.City_Category ## this the target
features = ['Occupation', 'Stay_In_Current_City_Years', 'Marital_Status', 'Purchase',
            'Product_Category_1', 'Product_Category_2', 'Product_Category_3']
X = blackFridayClassification2[features]

```

```

In [103]: ##### Converting the dataset into a Dmatrix will allow us to take advantage of the parallel processing
dmatrix = xgb.DMatrix(data=X, label=mytarget)

```

```

In [104]: ## We will take 70 % for training and 30 % for testing
x_train, x_test, y_train, y_test = train_test_split(X, mytarget, test_size=0.3, random_state=42)

```

Building and training Model

```

In [105]: model = XGBClassifier()
model.fit(x_train, y_train)

```

```

Out[105]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, gamma=0, learning_rate=0.1, max_delta_step=0,
                        max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
                        n_jobs=1, nthread=None, objective='multi:softprob', random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                        silent=True, subsample=1)

```

Making prediction with XGBoost Classifier

```

In [106]: y_pred = model.predict(x_test)
predictions = [round(value) for value in y_pred]

```

Test the performance of XGBoost

```

In [107]: accuracy = accuracy_score(y_test, predictions)
print(accuracy)

```

0.4488076193310763

The accuracy of the model is 44.88 % the model improves slightly with XGboost

k-fold Cross Validation using XGBoost

"In order to build more robust models, it is common to do a k-fold cross validation

where all the entries in the original training dataset are used for both training as well as validation. " Datacamp

```

In [108]: params = {"objective": "multi:softmax", 'num_class': 3, 'colsample_bytree': 0.3, 'learning_rate': 0.1,
                  'max_depth': 8, 'alpha': 10}

```

```

cv_results = xgb.cv(dtrain=dmatrix, params=params, nfold=3,
                    num_boost_round=100, early_stopping_rounds=10, metrics="mlogloss")

```

```

In [110]: cv_results.head()

```

```

Out[110]:

```

	train-mlogloss-mean	train-mlogloss-std	test-mlogloss-mean	test-mlogloss-std
0	1.095657	0.000023	1.095684	0.000008
1	1.091976	0.000058	1.092066	0.000034
2	1.087536	0.000616	1.087670	0.000570
3	1.084858	0.000849	1.085052	0.000778
4	1.082032	0.001115	1.082295	0.001123

cv_results contains train and test mlogloss for each boosting round. we take only the first 5 elements

Visualizing XGBoost Classifier

```

In [98]: xg_classifier = xgb.train(params=params, dtrain=dmatrix, num_boost_round=15)

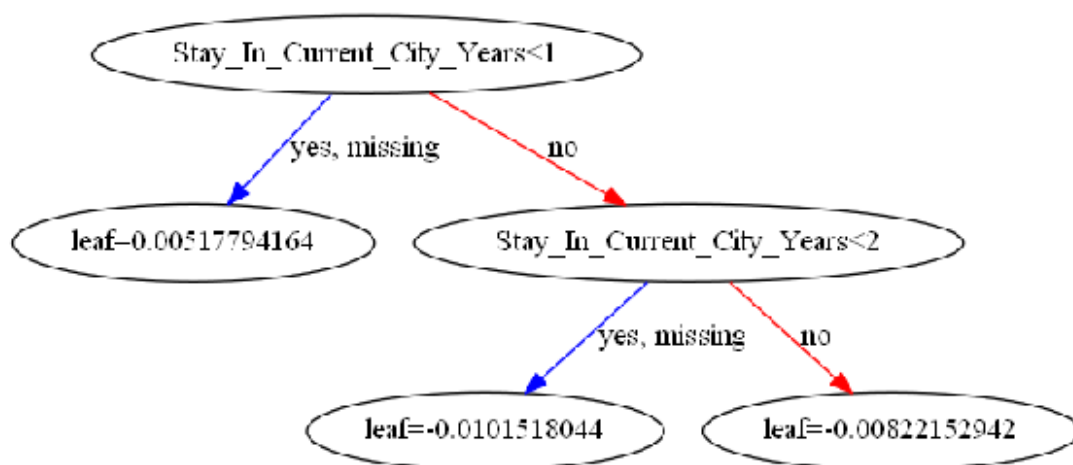
```



```

In [107]: xgb.plot_tree(xg_claasifier,num_trees=30)
plt.show()

```



we can conclude that XGboost classifier did improve the model

Algorithm #3 DBSCAN

We are going to use cluster analysis to segment customers based on their spending

this part is bit different fom the previous one because we do not have target

```

In [64]: ## we import the libraries require for running DBSCAN
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler
from sklearn import metrics
from sklearn.datasets.samples_generator import make_blobs

```

```

In [43]: ## import library for Principal component analysis
from sklearn.decomposition import PCA

```

```

In [54]: ## we select features that we wish to use in the clustering
clusterFeatures = ['Product_Category_1','Product_Category_2','Product_Category_3',

```

```

In [55]: ## we plot in X the features and in y Gender
X=blackFridayClustering[clusterFeatures]
y=blackFridayClustering.Gender

```

```
► In [47]: ## we scale the data  
X= StandardScaler().fit_transform(X)
```

```
c:\users\diall\appdata\local\programs\python\python37\lib\sklearn\preprocessin  
g\data.py:645: DataConversionWarning: Data with input dtype int64 were all con  
verted to float64 by StandardScaler.  
    return self.partial_fit(X, y)  
c:\users\diall\appdata\local\programs\python\python37\lib\sklearn\base.py:464:  
DataConversionWarning: Data with input dtype int64 were all converted to float  
64 by StandardScaler.  
    return self.fit(X, **fit_params).transform(X)
```

We are going to reduce the dimension from 4 to 2 dimensions using PCA

```
► In [48]: ## create the object PCA and fit the model  
pca = PCA(2) # project from 4 to 2 dimensions  
principalComponents = pca.fit_transform(X)  
principalDf = pd.DataFrame(data = principalComponents  
    , columns = ['Principal Component 1', 'Principal Component 2'])
```


```
► In [49]: # we combine the two component with the target that we are using which gender in t  
finalDf = pd.concat([principalDf, blackFridayClustering[['Gender']]], axis = 1)
```

This section is just plotting 2 dimensional data.

```

In [51]: fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel('Principal Component 1', fontsize = 15)
ax.set_ylabel('Principal Component 2', fontsize = 15)
ax.set_title('2 component PCA', fontsize = 20)
targets = ['F', 'M']
colors = ['r', 'b']
for target, color in zip(targets, colors):
    indicesToKeep = finalDf['Gender'] == target
    ax.scatter(finalDf.loc[indicesToKeep, 'Principal Component 1'],
              finalDf.loc[indicesToKeep, 'Principal Component 2'],
              c = color,
              s = 50)
ax.legend(targets)
ax.grid()

```



what we see from this result is that the class are not well separated

We are going to use Kmeans to see what the clustering will be

```
► In [45]: ## we reduce the sample size  
clustering=blackFriday.sample(frac=0.1, random_state=1)
```

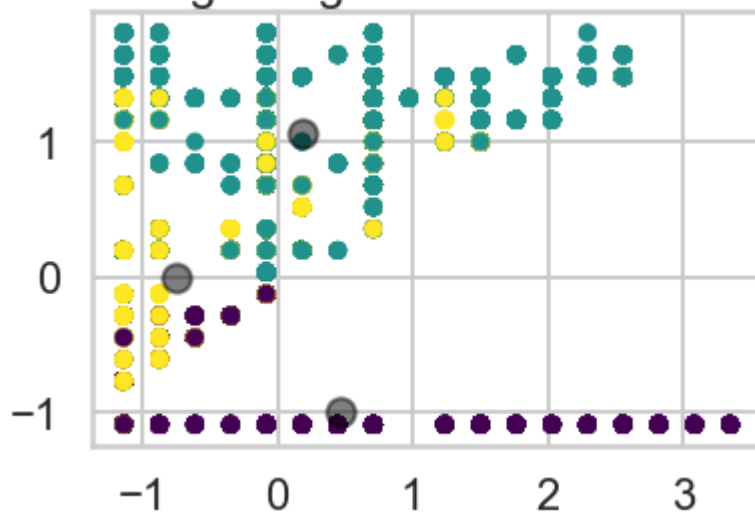
```
► In [74]: ## we import Kmeans libraries, create the model with 3 clusters, and fit the model  
from sklearn.cluster import KMeans  
kmeans = KMeans(n_clusters=3)  
kmeans.fit(X)  
y_kmeans = kmeans.predict(X)
```

We plot the result of the clustering

```
► In [79]: plt.scatter(X[:, 0], X[:, 1], c=y_kmeans, s=50, cmap='viridis')  
centers = kmeans.cluster_centers_  
plt.scatter(centers[:, 0], centers[:, 1], c='black', s=200, alpha=0.5);  
plt.title('Clustering using Kmeans with 3 clusters')
```

Out[79]: Text(0.5, 1.0, 'Clustering using Kmeans with 3 clusters')

Clustering using Kmeans with 3 clusters



we clearly see that the data are not around the cluster, that means the clustering did not work well

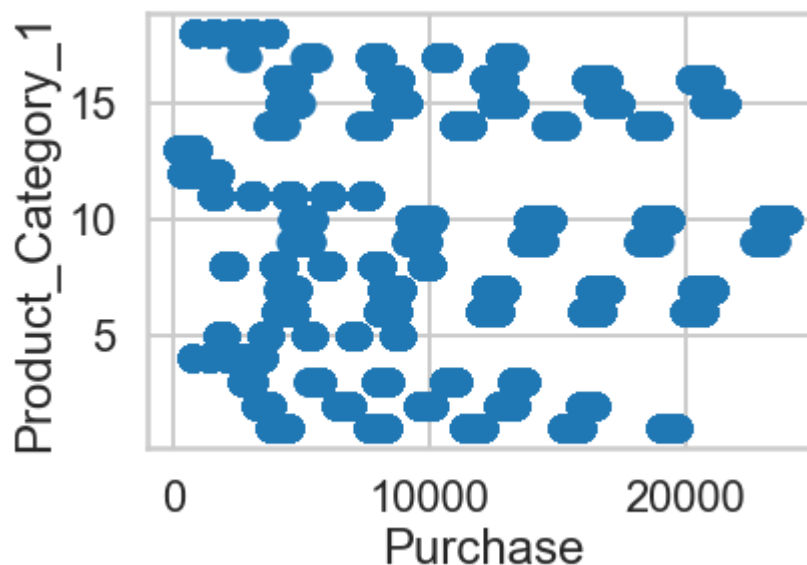
Let's plot the relationship

```

In [63]: x = blackFridayClustering['Purchase']
y = blackFridayClustering['Product_Category_1']

plt.scatter(x,y)    ## we plot the scatter
plt.xlabel("Purchase")
plt.ylabel("Product_Category_1")
plt.show()

```



We create the features for the DBSCAN model

```

In [65]: ## features selections
clusterFeatures = ['Product_Category_1', 'Product_Category_2', 'Product_Category_3', '

```

```

In [67]: features=clustering[clusterFeatures]

```

```

In [68]: ## we scale the data
stscaler = StandardScaler()

X = stscaler.fit_transform(features)

```

```

c:\users\diall\appdata\local\programs\python\python37\lib\sklearn\preprocessin
g\data.py:645: DataConversionWarning: Data with input dtype int64 were all con
verted to float64 by StandardScaler.

```

```

return self.partial_fit(X, y)
c:\users\diall\appdata\local\programs\python\python37\lib\sklearn\base.py:464:
DataConversionWarning: Data with input dtype int64 were all converted to float
64 by StandardScaler.
return self.fit(X, **fit_params).transform(X)

```

Building, training the model

```
► In [20]: db = DBSCAN(eps=0.3, min_samples=10).fit(X)
```

```
► In [69]: ## this line of code we select label  
labels = db.labels_  
core_samples = np.zeros_like(labels, dtype = bool)  
core_samples[db.core_sample_indices_] = True  
core_samples_mask = np.zeros_like(db.labels_, dtype=bool)  
core_samples_mask[db.core_sample_indices_] = True
```

```
► In [70]: # Number of clusters in labels, ignoring noise if present.  
n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)  
n_noise_ = list(labels).count(-1)
```

```
► In [71]: print('Estimated number of clusters: %d' % n_clusters_)  
print('Estimated number of noise points: %d' % n_noise_)
```

```
Estimated number of clusters: 195  
Estimated number of noise points: 176
```

```


In [72]: unique_labels = set(labels)
colors = [plt.cm.Spectral(each)
           for each in np.linspace(0, 1, len(unique_labels))]
for k, col in zip(unique_labels, colors):
    if k == -1:
        # Black used for noise.
        col = [0, 0, 0, 1]

    class_member_mask = (labels == k)

    xy = X[class_member_mask & core_samples_mask]
    plt.plot(xy[:, 0], xy[:, 1], 'o', markerfacecolor=tuple(col),
             markeredgecolor='k', markersize=14)

    xy = X[class_member_mask & ~core_samples_mask]
    plt.plot(xy[:, 0], xy[:, 1], 'o', markerfacecolor=tuple(col),
             markeredgecolor='k', markersize=6)
plt.title('Estimated number of clusters: %d' % n_clusters_)
plt.show()

```



we see from this result that the classes are not well separated

Test of performance of DBSCAN

```

In [73]: ## the result of the accuracy
score=metrics.silhouette_score(X,labels)
print('Percentage of accuracy: %f' % score)

```

Percentage of accuracy: 0.428055

The accuracy of the model is 42.80 %

this is the end of my project, my goal was to use Regression, Classification and Clustering

to see how I can help advice the retail store about their business based on data driven

References

<https://www.ritchieng.com/pandas-changing-datatype/>
(<https://www.ritchieng.com/pandas-changing-datatype/>)

<https://seaborn.pydata.org/generated/seaborn.boxplot.html>
(<https://seaborn.pydata.org/generated/seaborn.boxplot.html>)

<http://jonathansoma.com/lede/foundations-2017/classes/pandas-text-part-1/classwork/> (<http://jonathansoma.com/lede/foundations-2017/classes/pandas-text-part-1/classwork/>)

<https://towardsdatascience.com/analyze-the-data-through-data-visualization-using-seaborn-255e1cd3948e>
(<https://towardsdatascience.com/analyze-the-data-through-data-visualization-using-seaborn-255e1cd3948e>)

<https://www.kaggle.com/mburakergenc/predictions-with-xgboost-and-linear-regression> (<https://www.kaggle.com/mburakergenc/predictions-with-xgboost-and-linear-regression>)

https://www.w3cschool.cn/doc_scikit_learn/scikit_learn-auto_examples-cluster-plot_dbscan.html?lang=en
(https://www.w3cschool.cn/doc_scikit_learn/scikit_learn-auto_examples-cluster-plot_dbscan.html?lang=en)